

PostgreSQL Notes

by [Aryan](#)

Database Fundamentals

What is a Database?

A structured system for storing, managing, and retrieving data efficiently. 💾

PostgreSQL Essentials

PostgreSQL Overview

PostgreSQL, or Postgres, is a powerful open-source relational database system known for advanced features and modern design. 🚀

psql Command-Line

- `\?` - Help
- `\l` - List databases
- `\psql --help` - `psql help`

Connecting to a Database

Two ways:

1. `psql -h localhost -p 5432 -U username database_name`
2. Within `psql` shell: `\l`, `\c database_name`

Danger Zone

```
SOL
1 | DROP DATABASE name
```

⚠ Irreversible! Deletes all data and tables. Never use on production databases.

Creating Tables

Table Creation

Use `CREATE TABLE`. Define columns, data types, and constraints.

```
CREATE TABLE person (
  id BIGSERIAL NOT NULL PRIMARY KEY,
  first_name VARCHAR(50) NOT NULL,
  last_name VARCHAR(50) NOT NULL,
  gender VARCHAR(6) NOT NULL,
```

```
date_of_birth DATE NOT NULL
);
```

Describing Tables

\d or \d table_name to describe tables.

Data Manipulation

Inserting Data

Use INSERT INTO for data insertion:

```
INSERT INTO person (first_name, last_name, gender, date_of_birth)
VALUES
  (1, 'John', 'Doe', 'MALE', DATE '1990-05-15'),
  (2, 'Jane', 'Smith', 'FEMALE', DATE '1988-03-22'),
  (3, 'Michael', 'Johnson', 'MALE', DATE '1995-09-10'),
  (4, 'Emily', 'Brown', 'FEMALE', DATE '1992-07-03'),
  (5, 'David', 'Lee', 'MALE', DATE '1987-11-28'),
  (6, 'Olivia', 'Garcia', 'FEMALE', DATE '1998-02-19'),
  (7, 'Daniel', 'Martinez', 'MALE', DATE '1986-12-14'),
  (8, 'Sophia', 'Anderson', 'FEMALE', DATE '1991-04-27'),
  (9, 'Ethan', 'Taylor', 'MALE', DATE '1996-06-05'),
  (10, 'Ava', 'Clark', 'FEMALE', DATE '1989-08-08');
```

Querying Data

Selecting Data

Use SELECT :

```
SELECT * FROM person;
SELECT first_name, last_name FROM person;
```

Output:

id	first_name	last_name	gender	date_of_birth
1	John	Doe	MALE	1990-05-15
2	Jane	Smith	FEMALE	1988-03-22
3	Michael	Johnson	MALE	1995-09-10
4	Emily	Brown	FEMALE	1992-07-03
5	David	Lee	MALE	1987-11-28
6	Olivia	Garcia	FEMALE	1998-02-19
7	Daniel	Martinez	MALE	1986-12-14

id	first_name	last_name	gender	date_of_birth
8	Sophia	Anderson	FEMALE	1991-04-27
9	Ethan	Taylor	MALE	1996-06-05
10	Ava	Clark	FEMALE	1989-08-08

Sorting Data

- Ascending: `SELECT * FROM person ORDER BY last_name`
- Descending: `SELECT * FROM person ORDER BY last_name DESC`

DATA TYPES

Data Type	Description
<code>bigint</code>	Signed eight-byte integer
<code>integer</code>	Signed four-byte integer
<code>smallint</code>	Signed two-byte integer
<code>numeric(p,s)</code>	Exact numeric with <code>p</code> digits, <code>s</code> of them
<code>real</code>	Four-byte floating-point
<code>double precision</code>	Eight-byte floating-point
<code>decimal(p,s)</code>	Exact numeric with <code>p</code> digits, <code>s</code> of them
<code>text</code>	Variable-length character string
<code>char(n)</code>	Fixed-length character string
<code>varchar(n)</code>	Variable-length character string
<code>date</code>	Date (year, month, day)
<code>time</code>	Time of day (no time zone)
<code>timestamp</code>	Date and time (no time zone)
<code>timestamptz</code>	Date and time with time zone
<code>interval</code>	Time interval
<code>boolean</code>	True/False values
<code>enum</code>	Enumerated (enum) types
<code>uuid</code>	Universally unique identifier
<code>json</code>	JSON data type
<code>jsonb</code>	Binary JSON data type
<code>bytea</code>	Binary data
<code>point</code>	Point on a plane
<code>line</code>	Infinite line on a plane
<code>lseg</code>	Line segment on a plane
<code>box</code>	Rectangular box on a plane
<code>path</code>	Closed path on a plane
<code>polygon</code>	Polygon on a plane

Data Type	Description
circle	Circle on a plane
inet	IPv4 or IPv6 host address
cidr	IPv4 or IPv6 network address
macaddr	MAC address
bit(n)	Fixed-length bit string
bit varying(n)	Variable-length bit string
tsvector	Text search vector (full-text search)
tsquery	Text search query (full-text search)
hstore	Key-value store
oid	Object identifier (system use)
pg_lsn	Log sequence number
uuid	Universally unique identifier
money	Currency amount
xml	XML data
interval	Time interval
oidvector	Array of object identifiers

Distinct Values

```
SELECT DISTINCT gender FROM person;
```

Output:

gender
MALE
FEMALE

Filtering Data

```
SELECT * FROM person WHERE gender = 'FEMALE';
```

Output:

id	first_name	last_name	gender	date_of_birth
2	Jane	Smith	FEMALE	1988-03-22
4	Emily	Brown	FEMALE	1992-07-03
6	Olivia	Garcia	FEMALE	1998-02-19
8	Sophia	Anderson	FEMALE	1991-04-27
10	Ava	Clark	FEMALE	1989-08-08

Comparison Operators

Use `<`, `>`, `=`, `<=`, and `>=`.

Limiting Results

```
SELECT * FROM person LIMIT 5;
SELECT * FROM person OFFSET 5 LIMIT 5;
SELECT * FROM person OFFSET 5 FETCH 5 ROW ONLY;
```

Advanced Queries

IN Clause

Filter with `IN`:

```
SELECT * FROM person WHERE last_name IN ('Doe', 'Smith', 'Anderson');
```

Output:

id	first_name	last_name	gender	date_of_birth
1	John	Doe	MALE	1990-05-15
2	Jane	Smith	FEMALE	1988-03-22
8	Sophia	Anderson	FEMALE	1991-04-27

BETWEEN Clause

```
SELECT * FROM person
WHERE date_of_birth BETWEEN DATE '1990-01-01' AND '1995-12-31';
```

Output:

id	first_name	last_name	gender	date_of_birth
1	John	Doe	MALE	1990-05-15
4	Emily	Brown	FEMALE	1992-07-03
8	Sophia	Anderson	FEMALE	1991-04-27

LIKE Operator

The `LIKE` operator allows you to perform pattern matching within strings.

- `_` matches any single character.
- `%p%` matches any string containing 'p'.
- `ILIKE` can be used for case-insensitive matching.

Examples:

- Match names starting with 'J':

```
SELECT * FROM person WHERE first_name LIKE 'J%';
```

Output:

id	first_name	last_name	gender	date_of_birth
1	John	Doe	MALE	1990-05-15
2	Jane	Smith		

| FEMALE | 1988-03-22 |

- Match names containing 'an':

```
SELECT * FROM person WHERE last_name LIKE '%an%';
```

- Case-insensitive match for 'john':

```
SELECT * FROM person WHERE first_name ILIKE 'john';
```

GROUP BY Clause

The `GROUP BY` clause is used to group rows with similar values in specified columns, often used with aggregate functions.

Example:

```
SELECT country_of_birth, COUNT(*) FROM person GROUP BY country_of_birth;
```

Output:

country_of_birth	count
USA	3
Canada	2
UK	2
France	1
Germany	2

GROUP BY HAVING Clause

The `GROUP BY` clause can be combined with `HAVING` to filter grouped results based on aggregate functions.

Example:

```
SELECT country_of_birth, COUNT(*) FROM person GROUP BY country_of_birth HAVING COUNT(*) > 1;
```

Output:

country_of_birth	count
USA	3
Canada	2
UK	2
Germany	2

Useful Aggregate Functions

Sample Car Table 🚗

Let's create a sample `car` table with 10 values:

id	make	model	price
1	Ford	Mustang	50000
2	Honda	Civic	25000
3	Ford	Fusion	30000
4	Toyota	Camry	28000
5	Honda	Accord	32000
6	Toyota	Corolla	22000
7	Ford	Focus	27000
8	Honda	Fit	18000
9	Toyota	RAV4	35000
10	Ford	Escape	31000

MAX

To find the maximum value in a column:

```
SELECT MAX(price) FROM car;
```

Output:

```
max
-----
50000.00
(1 row)
```

To find the maximum value per category (e.g., make):

```
SELECT make, MAX(price) FROM car GROUP BY make;
```

Output:

```
make | max
-----|-----
```

Ford		50000.00
Honda		32000.00
Toyota		35000.00

(3 rows)

MIN

To find the minimum value in a column:

```
SELECT MIN(price) FROM car;
```

Output:

min

18000.00

(1 row)

To find the minimum value per category (e.g., make):

```
SELECT make, MIN(price) FROM car GROUP BY make;
```

Output:

make		min
-----		-----
Ford		27000.00
Honda		18000.00
Toyota		22000.00

(3 rows)

AVERAGE

To calculate the average value in a column:

```
SELECT AVG(price) FROM car;
```

Output:

avg

29500.0000000000

(1 row)

To round the average value:

```
SELECT ROUND(AVG(price)) FROM car;
```


Output:

```
round
-----
29500
(1 row)
```

SUM

To find the total sum of values in a column:

```
SELECT SUM(price) FROM car;
```

Output:

```
sum
-----
295000.00
(1 row)
```

To find the total sum per category (e.g., make):

```
SELECT make, SUM(price) FROM car GROUP BY make;
```

Output:

```
make | sum
-----|-----
Ford  | 138000.00
Honda | 75000.00
Toyota | 85000.00
(3 rows)
```

Arithmetic Operators

Arithmetic operators allow you to perform calculations on columns in your queries.

- For applying a percentage discount:

```
SELECT id, make, model, price, price * 0.10 AS discount FROM car;
```

- To round the calculated discount:

```
SELECT id, make, model, price, ROUND(price * 0.10, 2) AS discount FROM car;
```

- Calculating discounted price:

```
SELECT id, make, model, price, (price - (price * 0.10)) AS discounted_price FROM car;
```

Alias

You can use the `AS` keyword to assign aliases to columns or expressions in your query results.

```
SELECT id, make, model, price, price * 0.10 AS discount FROM car;
```

In this example, `price * 0.10` is given the alias "discount."

COALESCE 🙌

The `COALESCE` function returns the first non-null argument.

- To provide a default value if a column is null:

```
SELECT COALESCE(email, 'email not provided') FROM person;
```

In this example, if the `email` column is null, it returns 'email not provided.'

NULLIF 🚫

The `NULLIF` function compares two expressions and returns null if they are equal; otherwise, it returns the first expression.

- To avoid division by zero:

```
SELECT NULLIF(10 / 0, 10);
```

Here, it returns null to avoid division by zero.

- Combining `COALESCE` and `NULLIF`:

```
SELECT COALESCE(10 / NULLIF(0, 0), 0);
```

Date and Time

You can work with date and time values in PostgreSQL.

- To get the current date and time:

```
SELECT NOW();
```

- To extract specific components from dates:

```
SELECT EXTRACT(YEAR FROM NOW());  
SELECT EXTRACT(MONTH FROM NOW());  
SELECT EXTRACT(DAY FROM NOW());  
SELECT EXTRACT(DOW FROM NOW()); -- Day of the week  
SELECT EXTRACT(CENTURY FROM NOW());
```

- To calculate age based on date of birth:

```
SELECT AGE(NOW(), DOB) FROM person;
```

Primary Key

A primary key is a column or set of columns that uniquely identifies each row in a table. It enforces the uniqueness of values and ensures that the column(s) cannot contain null values.

```
ALTER TABLE person ADD PRIMARY KEY (id);
```

UNIQUE Constraint

A UNIQUE constraint ensures that all values in a column or a group of columns are distinct. It enforces the uniqueness of values but allows null values.

```
ALTER TABLE person ADD CONSTRAINT unique_email UNIQUE (email);
```

Foreign Key

A foreign key is a column or a set of columns in a table that is used to establish and enforce a link between the data in two tables. It creates a relationship between the tables.

```
ALTER TABLE orders
ADD CONSTRAINT fk_customer_id
FOREIGN KEY (customer_id)
REFERENCES customers (id);
```

Indexing

Indexes are database objects used to speed up data retrieval. They are created on columns to quickly locate and access data rows.

```
CREATE INDEX idx_last_name ON person (last_name);
```

Subqueries

A subquery is a query nested within another query. It can be used to retrieve data needed for the main query or for filtering results.

Example: Find all customers who have placed an order:

```
SELECT * FROM customers
WHERE id IN (SELECT DISTINCT customer_id FROM orders);
```

Joins

Joins are used to combine rows from two or more tables based on related columns between them.

INNER JOIN

An `INNER JOIN` returns only the rows that have matching values in both tables.

```
SELECT customers.name, orders.order_date
FROM customers
INNER JOIN orders ON customers.id = orders.customer_id;
```

LEFT JOIN (or LEFT OUTER JOIN)

A `LEFT JOIN` returns all rows from the left table and the matched rows from the right table. If there is no match, it returns NULL values for right table columns.

```
SELECT customers.name, orders.order_date
FROM customers
LEFT JOIN orders ON customers.id = orders.customer_id;
```

RIGHT JOIN (or RIGHT OUTER JOIN)

A `RIGHT JOIN` returns all rows from the right table and the matched rows from the left table. If there is no match, it returns NULL values for left table columns.

```
SELECT customers.name, orders.order_date
FROM customers
RIGHT JOIN orders ON customers.id = orders.customer_id;
```

FULL OUTER JOIN

A `FULL OUTER JOIN` returns all rows when there is a match in either the left or the right table. If there is no match, it returns NULL values for columns from the table without a match.

```
SELECT customers.name, orders.order_date
FROM customers
FULL OUTER JOIN orders ON customers.id = orders.customer_id;
```

Joins and Relationships

Creating Tables and Adding Relationship

```
CREATE TABLE people (
  id SERIAL PRIMARY KEY,
  first_name VARCHAR(50),
  last_name VARCHAR(50),
  gender VARCHAR(10),
  date_of_birth DATE
```

```
);

CREATE TABLE cars (
  id SERIAL PRIMARY KEY,
  make VARCHAR(50),
  model VARCHAR(50),
  price DECIMAL(10, 2),
  owner_id INT,
  FOREIGN KEY (owner_id) REFERENCES people(id),
  UNIQUE (owner_id, id)
);
```

Sample Data:

```
INSERT INTO people (first_name, last_name, gender, date_of_birth) VALUES
('Alice', 'Johnson', 'Female', '1990-05-15'),
('Bob', 'Smith', 'Male', '1985-02-10');
```

```
INSERT INTO cars (make, model, price, owner_id) VALUES
('Toyota', 'Camry', 25000.00, 1),
('Honda', 'Civic', 22000.00, 2),
('Ford', 'Mustang', 45000.00, 1);
```

Querying Data:

```
SELECT people.first_name, people.last_name, cars.make, cars.model
FROM people
JOIN cars ON people.id = cars.owner_id;
```

Output:

first_name	last_name	make	model
Alice	Johnson	Toyota	Camry
Bob	Smith	Honda	Civic
Alice	Johnson	Ford	Mustang

Handling Constraints and Operations ⚠️

DELETE Data 🗑️

The DELETE statement removes rows:

```
DELETE FROM person WHERE id = 1;
```

UPDATE Data 🔄

The UPDATE statement modifies data:

```
UPDATE person SET gender = 'Other' WHERE gender = 'Unknown';
```

INSERT Data

INSERT adds new rows:

```
INSERT INTO person (name, age) VALUES ('John', 30);
```

Upsert and Excluded

An upsert operation combines INSERT and UPDATE:




```
INSERT INTO person (id, name) VALUES (1, 'Alice')  
ON CONFLICT (id) DO UPDATE SET name = EXCLUDED.name;
```

Exporting to CSV

To export data to a CSV file:

```
\COPY (SELECT * FROM people) TO 'people.csv' WITH CSV HEADER;
```

Additional PostgreSQL Features

- Serial Sequences 
- PostgreSQL Extensions 
- PL/v8 
- UUIDs as Primary Keys 